

End-to-end Mechanised Proof of an eBPF Virtual Machine for Microcontrollers

Anonymized submission to CAV'22

Abstract. RIOT is a micro-kernel dedicated to IoT applications that adopts eBPF (extended Berkeley Packet Filters) to implement so-called femto-containers: as micro-controllers rarely feature hardware memory protection, the isolation of eBPF virtual machines (VM) is critical to ensure system integrity against potentially malicious programs. This paper proposes a methodology to directly derive the verified C implementation of an eBPF virtual machine from a Gallina specification within the Coq proof assistant. Leveraging the formal semantics of the CompCert C compiler, we obtain an end-to-end theorem stating that the C code of our VM inherits its safety and security properties of its Gallina specification. Our refinement methodology ensures that the isolation property of the specification holds in the verified C implementation. Preliminary experiments demonstrate satisfying performance.

Keywords: proof methodology · virtual machines · fault isolation.

1 Introduction

Hardware-enforced memory isolation is often not available on micro-controller units (MCU), as memory protection units usually trade coarse-grain isolation for significant price and performance overheads, and architecture dependencies: *Trustzone*, Sanctum [5], Sancus [28], *etc.* To mitigate development variability and cost, common practice for MCU operating system design (RIOT [3], *FreeRTOS*, *TinyOS*, *Fushia*, and others [12]) advises to run all the device's code stack in a shared memory space, which can only be reasonably safe if that code can be trusted. While standard in safety-critical system design, such a trust requirement is oftentimes unsuitable for networked MCUs, where the extensibility of the OS kernel at runtime is an essential functionality. When system reconfiguration does not affect the entire network (via, *e.g.* leader election), extensibility can easily be provided offline, by employing library OSs or unikernels [22], to reconfigure network endpoints independently (*e.g.* cloud apps). Otherwise, the best solution is to load and execute system extensions (configurations, protocols, firewalls, *etc.*) as assembly-level scripts like Wasm [11] or the Berkeley packet filters (BPF [23]), using an interpreter or a just-in-time (JIT) compiler on the target device.

Femto-containers To this end, RIOT adopts the extended Berkeley packet filters (eBPF) and tailors it to resource-constrained MCUs by implementing so-called femto-containers: tiny virtual machine instances interpreting eBPF scripts. Compared to more expressive languages, like Wasm, experiments show that RIOT's eBPF implementation: rBPF, requires a smaller memory footprint [38]. Instead, the Linux kernel features an eBPF JIT compiler whose security depends on a

sophisticated online verifier [27]. As an MCU architecture cannot host such a large verifier, executing JIT code would here imply delegation of trust to a 3rd party, offline, verifier. Nonetheless, and while a VM is of slower execution speed than JIT-compiled code, it can, however, run untrusted, erroneous, adversary code in an open, unattended and possibly hostile environment, and still isolate faults to the scope of the offending script.

Approach & Goals This paper investigates an approach that trades high-performance on low-power devices for defensive programming and runtime footprint. Our primary goal is to prevent faults that could compromise host devices and, by extension, force networked devices to reboot and resynchronize (*i.e.* fault tolerance protocols). To maximize trust in the implementation of rBPF, we investigate a methodology allowing the verified extraction of C code directly from its mechanically proved definition in Gallina, the functional language embedded in the proof assistant Coq[4].

Method To mechanically prove the correctness of an interpreter, a conventional approach consists in defining the reference semantics in a proof assistant and in showing that an executable optimized interpreter produces the same output. In this paper, our goal is to verify the interpreter of the virtual rBPF instruction set, implemented with the system programming language C. To this aim, we introduce a direct, end-to-end, validation workflow. The semantics of the source instruction set is directly defined by monadic functional terms in our proof assistant. We prove that this semantics enforces safety and security requirements regarding memory and branching. Then, C code is automatically derived from these monadic functional terms to implement the expected virtual machine. We prove that the extracted C code has the same stateful behavior as that of the corresponding source instructions, as defined by their monadic specification. Our method uses a monadic subset of Gallina of sufficient expressivity to specify rBPF’s semantics, supports the verified extraction of equivalent Clight code, while provably implementing all required defensive runtime checks.

Plan The rest of the paper is organised as follows. Section 2 states our contributions. Section 3 provides some background on Berkeley packet filters and its variant, CompCert and the ∂x code extraction tool. Section 4 presents our methodology for formally refining a monadic Gallina program into a C program. Section 5 defines the proof model of our virtual machine: its semantics and consistency and isolation theorems. Section 6 refines the proof model of our femto-container into a synthesis model ready for code generation with CompCert. Section 7 proves the refinement between the synthesis and implementation models. Section 8 introduces the pre-processor of our virtual machine: a verified verifier. Section 9 case studies the performance of our generated VM implementation with respect to off-the-shelf RIOT femto-containers. Section 10 presents related works and Section 11 concludes.

The source code and proofs of our virtual machine, its generated code and benchmark data are available to the CAV’22 evaluation committee on an anonymized repository¹.

2 Contributions

This article presents the first end-to-end verification and synthesis of a full-scale, real-world, virtual machine for the BPF instruction set family: CertrBPF, an interpreter tailored to the hardware and resources constraints of MCU architectures running the RIOT operating system. Consequently, our secondary contribution is methodological, as it is (to our knowledge) an equally first full-scale case study of verified refinement and extraction of an executable C program directly from its model in a proof assistant (*i.e.* without adjoined specifications using a domain-specific formalism). We propose a workflow of proof-oriented programming using the functional language Gallina embedded in the proof assistant Coq. Our goal is to benefit from both the proof efficiency of verified programming (automation apart) and the exhaustivity of theorem proving.

Requirements Implementing a fault-isolating virtual machine for MCUs faces two major requirements. A first requirement is to embed the virtual machine inside the RIOT micro-kernel, hence, to minimize its code size and execution environment. A second requirement is to minimize the verification gap between the proof model and the running code. An obvious approach would be to use the existing Coq extraction mechanism to compile the Gallina model into OCaml. To this end, Coq extraction to OCaml should then be trusted and the OCaml runtime would become part of the trusted computing base (TCB), and would further need to be trimmed down to fit space requirements. Another approach would be to directly generate C code using, *e.g.* F*’s KreMLin compiler, which does however not yet provide a mechanised equivalence guarantee between source Low* programs and extracted C code.

Proof Methodology Our ambition is to minimize the verification gap and provide an end-to-end security proof linking our Gallina model to the extracted C code. Our intended TCB is hence restricted to the Coq type-checker, the C semantics of the CompCert compiler and a pretty-printer for the generated C abstract syntax tree (AST). To reach this goal, our starting point is a model of the rBPF semantics. It is written in Gallina, the functional language of the Coq proof assistant. We use this proof model to certify that all the memory accesses are valid and isolated to dedicated memory areas, thus ensuring isolation. From this proof model, we then derive a synthesis model of which we extract an executable version in Clight, that we finally prove to perform the same state transitions.

¹ <https://anonymous.4open.science/r/AnonymousCAV22-59DC>

A certified rBPF interpreter This paper introduces CertrBPF: the verified model and implementation of rBPF in Coq. We formalize the syntax and semantics of all rBPF instructions, implement a formal model of its interpreter (femto-container), complete the proof of critical properties of our model, and extract and verify CompCert C code from this formalization. This method allows us to obtain a fully verified virtual machine. Not only is the Gallina specification of the VM proved kernel- and memory-isolated using the proof assistant, but the direct interpretation of its intended semantics as CompCert C code is, itself, verified correct. This yields a fully verified binary program of maximum security and minimal TCB and memory footprint: CertrBPF, a memory-efficient, fully isolated, kernel-level virtual machine that isolates any runtime software fault using defensive code and does not necessitate offline verification.

Systems Integration & Micro-benchmarks We integrate CertrBPF as a drop-in replacement for the corresponding non-verified rBPF module in the RIOT operating system. We then comparatively evaluate the performance of CertrBPF integrated in RIOT, running on various 32-bit microcontroller architectures. Our benchmarks demonstrate that, in practice, CertrBPF’s memory footprint is on par, and incurs only a small execution slow-down, well worth its security gains.

3 Background

This section describes essential features of rBPF, of the CompCert compiler, and of the ∂x code generation tool, that are required by our refinement methodology.

BPF, eBPF and rBPF Originally, the purpose of Berkeley packet filters (BPF, [23]) was network packet filtering. The Linux community extended it to provide ways to run custom in-kernel VM code, hooked into various subsystems, for varieties of purposes beyond packet filtering [9]. eBPF was then ported to micro-controllers, yielding RIOT’s specification: rBPF [37]. Just as eBPF, rBPF is designed as a 64-bit register-based VM, using fixed-size 64-bit instructions and a reduced instruction set architecture (ISA). rBPF uses a fixed-sized stack (512 bytes) and defines no heap interaction, which limits the VM memory overhead in random-access memory (RAM). The rBPF specification, however, does not define special registers or interrupts for flow control, nor support virtual memory: the host device’s memory is accessed directly and only guarded using permissions.

The CompCert Verified Compiler. CompCert [16] is a C compiler that is both programmed and proved correct using the Coq proof assistant. The compiler is structured into compiler passes using several intermediate languages. Each intermediate language is equipped with a formal semantics and each pass is proved to preserve the observational behaviour of programs.

The Clight Intermediate Language Clight [18] is a pivotal language which condenses the essential features of C using a minimal syntax. The Verified Software Toolchain (VST) [2] verifies C programs at the Clight level that are obtained by the CLIGHTGEN tool. Though we do not reuse the proof infrastructure of VST, we are reusing CLIGHTGEN in order to get a Clight syntax from a C program.

CompCert Values and Memory Model [18,17] Values and memories are shared across all the intermediate languages of CompCert. The set of values val is defined as follows:

$$val \ni v ::= Vint(i) \mid Vlong(i) \mid Vptr(b, o) \mid Vundef \mid \dots$$

A value $v \in val$ can be a 32-bit integer $Vint(i)$; a 64-bit integer $Vlong(i)$, a pointer $Vptr(b, o)$ consisting of a block identifier b and an offset o , or the undefined value $Vundef$. The undefined value $Vundef$ represents an unspecified value and is not, strictly speaking, an undefined behaviour. Yet, as most of the C operators are strict in $Vundef$, and because branching over $Vundef$ or dereferencing $Vundef$ are undefined behaviours, our proofs will ensure the absence of $Vundef$. CompCert values also include floating-point numbers; they play no role in the current development. CompCert’s memory consists of a collection of separate arrays. Each array a has a fixed size determined at allocation time and is identified by an uninterpreted block $b \in block$. The memory provides an API for loading values from memory and storing values in memory. Operations are parameterised by a memory chunk k which specifies how many bytes should be written or read and how to interpret bytes as a value $v \in val$.

For instance, the memory chunk $Mint32$ specifies a 32-bit value and $Mint64$ a 64-bit value. The function $load\ k\ m\ b\ o$ takes a memory chunk k , a memory m , a block b and an offset o . Upon success, it returns a value v obtained from the memory by reading bytes from the block b starting at index o . Similarly, the function $store\ k\ m\ b\ o\ v$ takes a memory chunk k , a memory m , a block b , an offset o and a value v . Upon success, it returns an updated memory m' which is identical to m except that the block b contains the value v encoded into bytes according to the chunk k starting at offset o . The isolation properties offered by CompCert memory regions are worth mentioning: $load$ and $store$ operations fail (return $None$) for invalid offsets o and invalid permissions.

The ∂x tool ∂x emerged from the toolchain used to design and verify the Pip proto-kernel [13]. Its aim was to allow writing most of Pip’s source code in Gallina in a style as close to C as possible. ∂x extracts C code from a Gallina source program in the form of a CompCert C AST. The goal of ∂x is to provide C programmers with readily reviewable code and thus avoid misunderstanding between those working on C/assembly modules (that access hardware) and those working on Coq modules (the code and proofs). To achieve this, the language that ∂x can handle is only a (C-like) subset of Gallina. The functions that are to be converted to C rely on a monad to represent the side effects of the computation, such as modifications to the CPU state. But ∂x does not mandate a particular monad for code extraction.

∂x 's workflow ∂x proceeds in two steps. First, given a list of Gallina functions, or whole modules, it generates a simple intermediate representation (IR) for the subset of Gallina it can handle. The second step is to translate this IR into a CompCert C AST. Since Coq has no built-in reflection mechanism, the first step is written in Elpi [7], using the Coq-Elpi plugin [35]. That step can also process external functions (appearing as `extern` extracted C code) to support separate compilation with CompCert. In order to obtain an actual C file, ∂x also provides a small OCaml function that binds the extracted C AST to CompCert's C pretty-printer. Even though the ∂x language is a small subset of Gallina, it inherits much expressivity from the use of Coq types to manipulate values. For example, one can use bounded integers (*i.e.* the dependent pair of an integer with the proof that it is within some given range), that can be faithfully and efficiently represented as a single `int` in C. To this end, ∂x expects a configuration mapping Coq types to C.

∂x memory management A major design choice in the C-like subset of Gallina used by ∂x is memory management: its generated code executes without garbage collection. This affects the Coq types that can actually be used in ∂x : recursive inductive types, such as lists, cannot automatically be converted. However, this Gallina subset is particularly relevant to programs in which one wants to precisely control memory management and decide how to represent data structures in memory. This is typically the case of an operating system or, in our case, the rBPF virtual machine.

4 A generic method for end-to-end verification in Coq

This section gives an overview of our methodology to derive a verified C implementation from a Gallina specification. In the following sections, the methodology will be instantiated to derive the C implementation of a fault-isolating rBPF virtual machine and its verifier. The novelty of our approach is to provide an end-to-end correctness proof, within the Coq proof assistant, that reduces the hurdle of reasoning directly over the C code.

As shown in the Figure 1, the original rBPF C implementation is first formalized by a proof model in Gallina, and the verification of expected properties (*e.g.* safety) is performed by the Coq proof assistant. The Gallina specification is then refined into an optimized (and equivalent) synthesis model ready for extraction of C code. The refinement & optimization principle employed by our method consists of deriving a C-ready implementation, in Gallina, that is as close as possible to the expected target C code. This principle allows to 1) prove the native optimizations correct (the simulation theorem), 2) improve the performance of the extracted code and, 3) facilitate the extracted code review and validation with the system designers.

From the C-ready Gallina implementation, we leverage ∂x to automatically generate C code. Verifying this extraction consists of two steps: 1/ the generated C code is first parsed as a CompCert Clight model by the CLIGHTGEN tool of VST

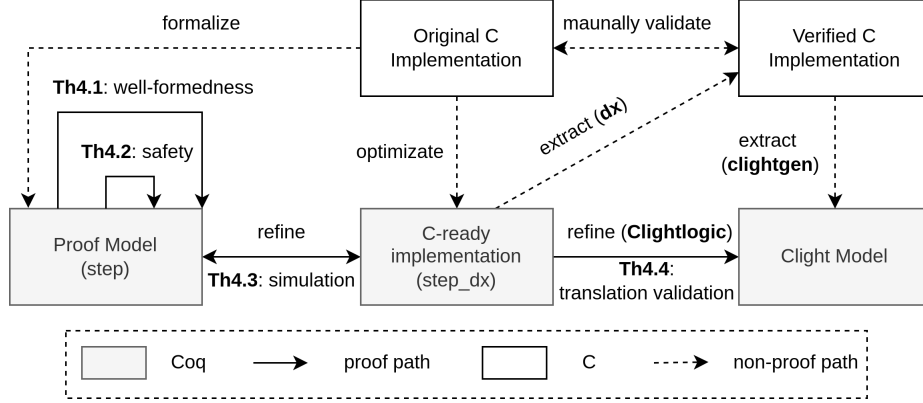


Fig. 1: End-to-end verification and synthesis workflow

and 2/ it is proved to refine the source Gallina model in Coq using translation validation.

Because ∂x generates C code in a syntax-directed manner, a minimal *Clight-logic* is designed to facilitate the refinement proof. The rest of the section explains these different steps in detail.

Proof-oriented specification. Our specification takes the form of an executable abstract machine in monadic form. It uses the standard option-state monad M .

$$\begin{aligned}
 M \text{ a state} &:= \text{state} \rightarrow \mathbf{option}(a \times \text{state}) \\
 \text{return}M &: a \rightarrow M \text{ a state} := \lambda st. \mathbf{Some}(a, st) \\
 \text{bind}M &: M \text{ a state} \rightarrow (a \rightarrow M \text{ b state}) \rightarrow M \text{ b state} := \\
 &\quad \lambda A.f.\lambda s.\mathbf{match} A \text{ s with } |\mathbf{None} \Rightarrow \mathbf{None} | \mathbf{Some}(x, s') \Rightarrow (f \ x) \ s'
 \end{aligned}$$

In the remainder, we write \emptyset for **None** and $[x]$ for **Some** x .

The monad threads the state along computations to model its in-place update. The safety property of the machine is implemented as an inline monitor: any violation leads to an unrecoverable error, *i.e.* the unique error represented by \emptyset . One step of the machine has the following signature:

$$\text{step} : M \ r \ \text{state}$$

where r is the type of the result. The *step* function implements a defensive semantics, checking the absence of error, dynamically. For our rBPF interpreter (see Sec. 5), the absence of error ensures that the rBPF code only performs valid instructions. In particular, all memory accesses are restricted to a sandbox specified as a list of memory regions. Function *step* is part of the TCB and, therefore, a mis-specification could result, after refinement, in an invalid computation. The purpose of the error state is to specify state transitions that would escape the scope of the safety property and, therefore, shall never be reachable from a well-formed state: $wf : \mathcal{P}(\text{state})$. We prove that well-formedness is an inductive property of the *step* function.

Theorem 1 (Well-formedness). *The step function preserves well-formedness.*

$$\forall st \ st'. \ st \in wf \wedge step \ st = \lfloor st' \rfloor \Rightarrow st' \in wf$$

We also prove that well-formedness is a sufficient condition to prevent the absence of error and, therefore, the safety of computations.

Theorem 2 (Safety). *The step function is safe, i.e. a well-formed state never leads to an error.*

$$\forall st. \ st \in wf \Rightarrow step \ st \neq \emptyset$$

C-ready implementation Our methodology consists in refining the *step* function into an interpreter $step_{\partial x}$ complying with the requirements of ∂x . As ∂x performs syntax-directed code generation, the efficiency of the extracted code crucially depends on $step_{\partial x}$. In order to preserve the absence of errors, we prove a simulation relation between the *step* and $step_{\partial x}$ functions. A direct consequence of the simulation theorem is that $step_{\partial x}$ never raises an error.

Theorem 3 (Simulation). *Given simulation relations $R_s \subseteq state \times state'$ and $R_r \subseteq r \times r'$, the function $step_{\partial x}$ simulates the function *step*.*

$$\forall s_1, s'_1, s_2, r. (s_1, s_2) \in R_s \wedge step \ s_1 = \lfloor r, s'_1 \rfloor \Rightarrow \exists s'_2, r'. \bigwedge \begin{cases} step_{\partial x} \ s_2 = \lfloor r', s'_2 \rfloor \\ (s'_1, s'_2) \in R_s \\ (r, r') \in R_r \end{cases}$$

Translation Validation of C code The next stage consists in refining the $step_{\partial x}$ function into a Clight program by relying on ∂x to get a C program and on the CLIGHTGEN tool to get a Clight $step_C$ program (see Sec. 6). As this pass is not trusted, we prove the following translation validation theorem.

Theorem 4 (Translation Validation). *Given a simulation relation $R_s \subseteq state' \times val \times mem$ and a relation $R_r \subseteq res \times val$, the Clight code $step_C$ refines the function $step_{\partial x}$:*

$$\begin{aligned} \forall s, v, k, m. (s, v, m) \in R_s \Rightarrow step_{\partial x} \ s = \lfloor (r, s') \rfloor \Rightarrow \\ \exists m', r'. Callstate(step_C, [v], k, m) \rightarrow^{*t} ReturnState(r', call_cont(k), m') \wedge \\ (s', v, m') \in R_s \wedge (r, r') \in R_r \end{aligned}$$

Theorem 4 states that, if $step_{\partial x} \ s$ runs without error and returns a result (r, s') , then, the Clight function $step_C$ successfully runs with argument v and, after a finite number of execution steps, returns a result r' and a memory m' that preserve the refinement relations. In our encoding, the unique argument v is a pointer to the memory allocated region refining the interpreter state and k represents the continuation of the computation. A corollary of Theorem 4 is that the Clight code $step_C$ is free of undefined behaviours. In particular, all memory accesses are valid. As the memory model does not allow to forge pointers, this yields a strong isolation property. The remainder of this paper presents the implementation of all the definitions and theorems stated of this section using the Coq proof assistant.

5 A proof-oriented virtual machine model

The proof model of our memory-isolated VM first requires preliminary definitions to denote the syntax and state of the interpreter, and auxiliary functions, to denote dynamic security checks. The rBPF instruction set, Figure 2, features binary arithmetic and logic operations, negation, (un)conditional jumps relative to an offset, operations to load/store values from/to registers/memory, and termination. It operates on 64-bit *registers* $\{R0, \dots, R10\}$, categorized as 32-bit immediate or sources *src* and destination registers *dst* modulo 16-bit offsets *ofs*. The source operand of a **Load** instruction can only be a register *reg*. Compared to the RIOT rBPF implementation, our specification does not comprise the call instruction to a trusted system API. As its control flows outside of the VM, responsibility for its security is delegated to the host OS (after sanitation of the caller arguments).

(Operands)	$dst, reg \in registers, src \in registers \cup immediate, ofs \in offset$
(Chunk)	$chk ::= byte \mid halfword \mid word \mid doublewords$
(Operators)	$op ::= add \mid sub \mid mul \mid div \mid and \mid or \mid$ $lsh \mid rsh \mid mod \mid xor \mid mov \mid arsh$ $cmp ::= eq \mid neq \mid lt \mid gt \mid le \mid ge \mid set \mid slt \mid sgt \mid sle \mid sge$
(Instruction)	$ins ::= \mathbf{Exit} \mid \mathbf{Neg} \ dst \mid \mathbf{Ja} \ ofs \mid \mathbf{Jump} \ cmp \ dst \ src \ ofs$ $\mid \mathbf{Alu} \ op \ dst \ src \mid \mathbf{Load} \ chk \ dst \ reg \ ofs \mid \mathbf{Store} \ chk \ dst \ src \ ofs$

Fig. 2: Core syntax of rBPF instruction set

Machine state A semantic state st is a tuple $\langle I, L, R, F, M, MRs \rangle$ consisting of a sequence of instructions I , the current location L , registers R , an interpreter flag F , a memory M and memory regions MRs . The flag F indicates the state of the rBPF interpreter – normal BPF_OK , denoted F_n – terminated BPF_SUCC , denoted F_t – or error-tolerable $BPF_ILLEGAL_DIV/MEM, \dots$, denoted F_e .

A collection of memory regions MRs forms a CompCert memory model. A region $mr = \langle start, size, perm, ptr \rangle$ is defined by its physical *start* address, its *size*, its access *permission* and a pointer *ptr* ($= Vptr \ b \ 0$) to the corresponding block b where the CompCert memory model stores it. We write $I(L_{pc})$ for the instruction located at the program counter L_{pc} . $R[r]$ retrieves the value of the register r in the register map R . Functions *Alu* and *Cmp* reuse the CompCert’s operators over the *val* type. The *Alu* function returns **None** if an error occurs, *e.g.* division by zero. Functions *load* and *store* are those of CompCert’s memory model (see Sec. 3).

Alu : $op \rightarrow val \rightarrow val \rightarrow option \ val$ **Cmp** : $cmp \rightarrow val \rightarrow val \rightarrow bool$
load : $chk \rightarrow mem \rightarrow block \rightarrow Z \rightarrow option \ val$
store : $chk \rightarrow mem \rightarrow block \rightarrow Z \rightarrow val \rightarrow option \ mem$

Dynamic checks Function *check_alu* dynamically checks the validity of an arithmetic to avoid *div-by-zero* and *undefined-shift* errors: for division instructions, *check_alu* requires the value *v* to be non-zero. The value *v* of an arithmetic or logical shift instructions should be within the specific range $n \in \{32, 64\}$ of the given instruction. CertrBPF has both 32-bit and 64-bit ALU instructions, but this paper only regards the latter for simplicity.

$$check_alu(op, v) \stackrel{\text{def}}{=} \begin{cases} v \neq 0 & , \text{ if } op \in \{div, mod\} \\ 0 \leq v < n & , \text{ if } op \in \{lsh, rsh, arsh\} \\ true & , \text{ otherwise} \end{cases}$$

Function *check_mem* returns a valid pointer (*Vptr b ofs*) if there exists a unique memory region *mr* in *MRs* such that 1) the permission *mr.perm* is at least *Readable* for *Load* and *Writable* for *Store*; i.e. *mr.perm* $\geq p$; and 2) the offset *ofs* is aligned, i.e. $ofs \% Z(chk) = 0$ and in bounds, i.e. $ofs \leq max_unsigned - Z(chk)$, and the interval $[ofs, hi_ofs]$ is in the range of *mr*. Otherwise, *check_mem* returns the null pointer *Vnullptr*. The function *Z(chk)* maps *byte*, *halfword*, *word* and *double* to 1, 2, 4, and 8.

$check_mem(p, chk, addr, MRs) \stackrel{\text{def}}{=} \text{if } \exists! mr \in MRs, b.$
let *ofs* := *addr* − *mr.start* **and** *hi_ofs* := *ofs* + *Z(chk)* **in**
 (*mr.ptr* = *Vptr b 0*) \wedge (*mr.perm* $\geq p$) \wedge (*ofs* $\% Z(chk) = 0$) \wedge
 (*ofs* $\leq max_signed - Z(chk)$) \wedge ($0 \leq ofs \wedge hi_ofs < mr.size$)
then *Vptr b ofs* **else** *Vnullptr*

Semantics Functions *interp* and *sem* formalize the implementation of our proof model M_p in the Coq proof assistant by defining a monadic interpreter of rBPF. The top-level recursion *interp* processes a (monotonically decreasing) *fuel* argument and a sequence of instructions *I* in nominal state F_n . The second stage *sem* processes individual instructions $I(L_{pc})$. *MR* and *I* are read-only. During normal execution, the flag remains F_n . If the flag turns to F_t or F_e while processing an instruction, execution stops. For instance, if the *fuel* is used up, the flag turns to F_e . Result *None* marks transitions to crash states that the proof assistant should weave as unreachable (ghosts) using carefully crafted definitions of the *check_alu* and *check_mem* functions. Note that the interpreter *interp* does not check the range of branching offsets (i.e. $0 \leq s.L < \text{length}(s.I)$) and register-out-of-bounds which are to be statically verified by the verifying pre-processor, Sec. 8.

```
interp = λfuel s. if fuel == 0 then Some((), s{F=Fe}) else
  match sem s with
  | Some ((), t) => if t.F ≠ Fn then Some((), t)
                    else interp (fuel-1) t{L = t.L+1}
  | None => None

sem = λs. match s.I(s.L) with
| Exit => Some ((), s{F = Ft})
| Ja ofs => Some ((), s{L = s.L+ofs})
```

```

| Jump cmp dst ofs => if Cmp(cmp, s.R[dst], s.R[src])
    then Some ((), s{L = s.L+ofs})
    else Some ((), s)
| Neg dst => Some ((), s{R[dst] = ¬ s.R[dst]})
| Alu op dst src => if check_alu(op, s.R[src]) then
    match Val(op, s.R[dst], s.R[src]) with
    | Some v => Some ((), s{R[dst] = v})
    | None => None
    else Some ((), s{F = Fe})
| Load chk dst reg ofs =>
    match check_mem(Readable, chk, s.R[reg]+ofs, s.MRs) with
    | Vptr b ofs => match load(chk, s.M, b, ofs) with
        | Some v => Some ((), s{R[dst] = v})
        | None => None
    | _ => Some ((), s{F = Fe})
| Store chk dst src ofs =>
    match check_mem(Writable, chk, s.R[dst]+ofs, s.MRs) with
    | Vptr b ofs => match store(chk, s.M, b, ofs, s.R[src]) with
        | Some N => Some ((), s{M = N})
        | None => None
    | _ => Some ((), s{F = Fe})
| _ => Some ((), s{F = Fe})

```

Exit terminates the program with flag F_t and unconditional jump *Ja* increments the *pc* by *ofs*. A conditional *Jump* does so when function *Cmp* says that “*src cmp dst*” is true or skips. An arithmetic operation *Alu op dst src*, *check_alu* first checks the validity of *op* with source *src*, evaluates *op* against destination *dst* using *Alu*, stores the result *v* in register *dst*. For simplicity, we omit the case of immediate *srcs*, which would read *src* instead of $R[src]$ in the rules below. If the result is *None*, so becomes the monadic state (undefined behavior). Our definition of *check_alu*, and well-formedness conditions (see Sec. 5.1) ensures that this will never happen and that, in case of error, the execution terminates with flag F_e . The *Neg* instruction is a simpler *Alu* instruction with only one operand (*Neg*). Similarly, the semantics of memory instructions discerns legal accesses (*Load-Store*) from invalid ones using the *check_mem* function, whose careful definition should avoid undefined behaviors. We write $s.F$ for the value of field F in record s and $s\{F = v\}$ updates it to v .

5.1 Proof of software-fault isolation

Our proof model M_p formalizes the semantics of rBPF. It is implemented in Coq using Gallina. Assessing its correctness consists of proving two essential properties: 1) the well-formedness of the virtual machine’s state, that is, its registers, memory and verifier invariants, and 2) software-fault isolation, that is, the isolation of all transitions to a crash state \emptyset using runtime safety checks (e.g. *check_alu*, *check_mem*), ergo the impossibility of a transition to an undefined behavior.

The register invariant states that all registers contain 64-bit integer values. This rules out 32-bit integers, *Vundef* but also pointers and floating-point numbers, for which the **Alu** function would be undefined.

Definition 1 (registers_inv). $\forall r \in \text{register}. \exists l. R[r] = V\text{long } l$

The memory consistency invariant is, expectably, a bit more administrative. It states that each CompCert memory region *mr* register 8-bit integer blocks *b* of memory *m*, designated by a pointer *mr.ptr* to the 32-bit physical *mr.start* address of *b*, the 32-bit *mr.size* of *b* and at least *Readable* permissions *mr.perm* across $[0, \text{size})$. Finally, every two regions should point to disjoint physical address spaces in *m* (as per CompCert's memory regions for $mr'.ptr \neq mr.ptr$).

Definition 2 (memory_inv). $\forall mr \in MRs, m. \exists b, start, size. s.t.$
 $mr.ptr = Vptr \ b \ 0 \wedge Mem.valid_block \ m \ b \wedge is_byte_block \ b \ m \wedge$
 $mr.start = Vint \ start \wedge mr.size = Vint \ size \wedge mr.perm \geq Readable \wedge$
 $Mem.range_perm \ m \ b \ 0 \ (Int.unsigned \ size) \ Cur \ mr.perm \wedge$
 $(\forall mr' \in MRs, mr' \neq mr \rightarrow mr'.ptr \neq mr.ptr)$

Linux' eBPF requires a verifier to statically analysis input applications in order to heuristically reject invalid cases. Instead, CertrBPF's verifier, introduced in Sec. 8, is abstracted by the invariant 3 below. This invariant stipulates the minimal pre-condition for the interpreter to accept sequence of instructions *I*. It states that each $I[i]$ of these instructions should reference registers within $[0, 10]$ and that the offset *ofs* of all jump instructions should be in range, i.e. $0 \leq i + ofs + 1 \leq length(I) - 1$.

Definition 3 (verifier_inv). $\forall i, I, ofs. 0 \leq i \leq length(I) - 1 \rightarrow$
 $0 \leq get_dst(I[i]) \leq 10 \wedge 0 \leq get_src(I[i]) \leq 10 \wedge$
 $((I[i] = Ja \ ofs \vee I[i] = Jump \ \dots \ ofs) \rightarrow$
 $0 \leq i + ofs + 1 \leq length(I) - 1)$

These three invariants implement well-formedness as proposed in Sec. 4. Therefore, the following Coq Theorem **sem_preserve_inv** instantiates Theorem 1 and states that well-formedness is preserved by the *interp* function. Similarly, Theorem **inv_ensure_no_undef** instantiates Theorem 2. This proves that the dynamic checks of the model M_p are sufficient to ensure the absence of error. In particular, all memory accesses are valid and performed within the dedicated memory regions. As a result, our model ensures software fault isolation property.

```

Theorem sem_preserve_inv:  $\forall (st \ st': \text{state}) (fuel: \text{nat})$ 
  (Hin: registers_inv st  $\wedge$  memory_inv st  $\wedge$  verifier_inv st)
  (Hsem: interp fuel st = Some (tt, st')),
  registers_inv st'  $\wedge$  memory_inv st'  $\wedge$  verifier_inv st'.

Theorem inv_ensure_no_undef:  $\forall (st: \text{state}) (fuel: \text{nat})$ 
  (Hin: registers_inv st  $\wedge$  memory_inv st  $\wedge$  verifier_inv st),
  interp fuel st  $\neq$  None.

```

The corollary of Theorems `sem_preserve_inv` and `inv_ensure_no_undef` is that our virtual machine, obtained by refinement of the proof model, will always isolate code from other memory regions of the operating system and will never crash it.

6 A synthesis-oriented eBPF interpreter

The coding style of the proof model M_p is quite different from the original RIOT implementation in C and lacks optimizations used in the latter to improve runtime performance. To this aim, the synthesis model M_s firstly refines M_p into an optimized, safe and behaviorally equivalent monadic model which is then automatically transformed into an effectful implementation model M_c using ∂x .

Synthesis model M_s . In order to extract a verified implementation of rBPF of competing efficiency with RIOT's, the synthesis model refines our proof model in several respects.

1/ *Renaming.* Each Coq inductive type may correspond to several C types (e.g. *Vint/Vlong* to *signed* or *unsigned*, 32-bit or 64-bit) and require careful encoding. The case of *Vptr* is particularly delicate, as the target type contextually relies on bit-size and signedness. To sort this out, we rename Coq types to match the correct C type. For example, *val64.t*, *valu32.t*, *vals32.t* are *Val* types mapped to `unsigned long long`, `unsigned int` and `int`, respectively.

2/ *Erasure.* There are *ghost functions* present in M_p to identify undefined behaviors (branches to state \emptyset). These branches can be erased in M_s since our safety theorem proves that no undefined behaviors is possible.

3/ *Masking.* M_p implements an eager decoding of all 64-bit binary instructions and all specific fields. However, most instructions do not use all of these fields. Instead, M_s adopts a fine-grained decoding, *i.e.* masking operations, yielding a verifiable optimisation reducing computation and reusing code. For example, given a binary instruction '0f 02 00 00 00 00 00 00', M_p first computes *opcode* = 0x0f, *dst* = R0, *src* = R2, *imm* = ..., then compares the opcode with all 78 possible ones to find out that the instruction is *add64 dst src*. Instead, M_s first masks '*opcode*&0x07' to identify *ALU64 dst*, then masks '*opcode*&0x08' to match *ALU64 dst src*, compares '*opcode*&0xf0' with the 13 possible alu opcodes, and only computes *dst* and *src*.

Equivalence. Both M_p and M_s use the same monadic state *st* as in Sec. 5. Hence, the simulation relation $R \subseteq st \times st$, required by the **Simulation** theorem, simply is equality. rBPF's **Simulation** theorem is defined with *interp* : *nat* \rightarrow *M unit*, the M_p interpreter, and *interp_dx* : *nat* \rightarrow *M unit*, the M_s interpreter, in Coq.

```
Theorem equivalence_relation:  $\forall$  (st: state) (fuel: nat),
  interp fuel st = interp_dx fuel st.
```

∂x configuration \mathcal{E} Implementation model M_c . To extract the implementation model, we supply ∂x with our monad M and a mapping relation from Gallina to C, Table 1. Inductive types map to C types, *e.g.* *reg* to *unsigned int* (note that a many-to-one relation from Gallina to C is legal). Gallina constructs & constant functions map to C operators & constants, *e.g.* ‘*Val.addl*’ to ‘+’, ‘*Int.repr*(-2)’ and ‘*true*’ to ‘-2’ and ‘1’, etc. Gallina functions map to C functions. For any function operating the monadic state, the target C function holds an additional argument *st* of type *struct state** which corresponds to the implicit state of the monad. Gallina’s *match-pattern* translates to C’s *switch-case*, etc.

Table 1: Mapping relation from Gallina to C

	Gallina	C
Types	<i>reg/sint32_t/valptr8_t ...</i>	<i>unsigned int/int/unsigned char* ...</i>
Constructions	<i>true/Int.repr(-2)/BPF_OK/Mint32 ...</i>	<i>1/-2/0/4...</i>
Constants	<i>Val.addl/Val.subl/Val.mull/Z.eqb ...</i>	<i>+/-/*/= / ...</i>
Functions	<i>eval_pc: M sint32_t ...</i>	<i>int eval_pc(struct state *) ...</i>
Code struct	<i>if-then-else, match-pattern ...</i>	<i>if-else, switch-case ...</i>

Code extraction with ∂x . The extracted C implementation preserves the structure of the original Gallina code, and the extracted C functions directly operate on actual memory locations as CompCert memory operations map to C expressions with a dereference. Consider the example of the *step_mem_st_reg* function.

```

Definition step_mem_st_reg (src64: val64_t) (addr: valu32_t) (dst:
  ↪ reg) (op: int8_t): M unit :=
  do opcode_st <- get_opcode_mem_st_reg op;
  match opcode_st with
  | op_BPF_STXW =>
    do addr_ptr <- check_mem Writable Mint32 addr;
    if eq_ptr_null addr_ptr then upd_flag BPF_ILLEGAL_MEM
    else (** i.e. Mem.storev Mint32 addr_ptr src64 *)
      do _ <- store_mem_reg Mint32 addr_ptr src64; returnM tt ...

```

CompCert’s Byte *int8_t* and the register type *reg*, are mapped to *unsigned char* and *unsigned int*. Constructs *op_BPF_STXW*, *BPF_ILLEGAL_MEM* and *Writable* are respectively mapped to ‘99’, ‘-2’ and ‘2U’. The constant function *eq_ptr_null* is translated into an operation to check whether a pointer is null. The ‘*match opcode_st with*’ is extracted to ‘*switch (opcode_st) case*’. Functions *step_mem_st_reg*, *check_mem* and *store_mem_reg* in C have an additional monadic argument *st*.

```

void step_mem_st_reg(struct bpf_state* st, unsigned long long src64,
  ↪ unsigned int addr, unsigned int dst, unsigned char op){
  unsigned char opcode_st; unsigned char *addr_ptr;
  opcode_st = get_opcode_mem_st_reg(op);

```

```

switch (opcode_st) {
case 99:
  addr_ptr = check_mem(st, 2U, 4U, addr);
  if (addr_ptr == 0) { upd_flag(st, -2); return;
  } else { // i.e. *(unsigned int *) addr_ptr = src64
    store_mem_reg(st, 4U, addr_ptr, src64); return; } ...

```

7 Simulation Proof of the C rBPF Virtual Machine

In this section, we explain how to establish Theorem 4 for the Clight code of our virtual machine, derived from ∂x , and compiled into a Clight AST in Coq using the CLIGHTGEN tool.

Simulation Relation. A crucial ingredient of this theorem is the simulation relation between the Gallina state monad and the Clight state which is essentially made of a CompCert memory. The Gallina state comprises a CompCert memory that models the various memory regions available to the rBPF program. This memory may also contain other blocks that are not modified by the virtual machine but represent other kernel data-structures. The simulation relation stipulates that such blocks also exist in the Clight memory and have the same content. The Clight memory contains additional blocks (*i.e.* *state_block*, *ins_block* and *mrs_block*) to model the other fields of the Gallina state. The layout and content of those blocks is depicted in Figure 3.




Fig. 3: Simulation relation R between st_{rbpf} , left, and $rBPF\text{Clight}$, right.

Solid arrows in Figure 3 are simulation relations between $state_block$ and st_{rbpf} . Solid lines are the equalities between the rBPF memory m and blocks in rBPFCLight memory. Dashed lines indicate relations of pointers to blocks in CompCert memory. The encoding exploits the fact that each field of the Gallina state has a known length. Thus, every field can be encoded as a continuous sub-block. As a result, the program counter is obtained from the first 4 bytes: loading a memory chunk of type *Mint32* at offset 0 retrieves the *pc* field of the Gallina state. The next 4 bytes encode the enumerated type flag. Here, each constructor of type flag is assigned an integer. The next 10×64 bits are used to encode the register bank of the Gallina state.

$$Rs(state, state_block, m) \stackrel{\text{def}}{=} \begin{cases} st_{rbpf}.pc & = \text{load } Mint32 \ m_{clight} \ state_block \ 0 \\ st_{rbpf}.flag & = \text{load } Mint32 \ m_{clight} \ state_block \ 4 \\ st_{rbpf}.R0 & = \text{load } Mint64 \ m_{clight} \ state_block \ 8 \\ \dots & \end{cases}$$

The next elements of the Clight block represent the lists of instructions and of memory regions. In a functional language, lists are potentially of unbounded length and have a polymorphic type. Here, our lists always have fixed lengths and elements of fixed size. As a result, a list is directly encoded by a field specifying its length followed by a pointer to its memory block. The elements of the list are stored continuously in the pointed block.

Systematic Proof of Simulation. Since the ∂x tool is syntax-directed, there is a systematic correspondence between source Gallina and target C code. We exploit this property to design a minimal Clight logic geared toward our simulation proof. Our *Clightlogic* generalizes the translation validation theorem of Section 4 (Theorem 4) to accommodate Gallina functions and C functions with multiple arguments. In that case, we have a precondition which states that the Gallina and C arguments are linked pairwise by a refinement relation. Most of the arguments are numeric values and, in this case, the refinement relation states that the Gallina and C values are the same. The *Clightlogic* also provides a syntax-directed proof principle for each pair of Gallina/C syntactic construct. For instance, the *bindM* operator translates to a sequence in the C code. Also, the result of a Gallina function call is bound to a local variable in C, *e.g.* the local variable v below stands for the monadic state in C and points to the state memory block.

$$\partial x(bindM \ f \ (\lambda x.g)) = (vx = f_C(v); g_C(v, vx))$$

To exploit this pattern, our invariants take the form of an association list mapping each local variable to a set of C values that is obtained by partially evaluating a refinement relation with the Gallina value computed by the function (Figure 3). To evaluate f , one needs to have a refinement relation Rs between the Gallina state st and the C value of v in memory m . Now, suppose that $fst = \lfloor r, st' \rfloor$. Since f_C is a correct refinement of f , we have $Rs(st', v, m')$ and $Rr(r, x)$ for the value x of the local variable vx in the current environment. We conclude by mapping $vx \mapsto Rr \ r$ and use this invariant for the refinement of g by g_C .

The translation validation theorem proves a forward simulation relation from Coq to Clight. A backward simulation relation could easily be constructed as the Coq subset we use is *receptive* and Clight is *determinate*[33].

8 CertrBPF verifier

Linux eBPF's compiler and runtime system do not enforce type or memory safety. Instead, safety is verified prior to execution using a static analyser that checks programs validity (*e.g.* access to arbitrary kernel data structures). As both the size and complexity cannot fit the requirements of an MCU architecture, CertrBPF instead provides a simple (linear time) but formally verified analyzer to play a similar role of pre-processor for the interpreter. CertrBPF-verifier, aims at providing a minimal assumption of CertrBPF-interpretter to statically check the invariant *verifier_inv* (Def. 3). Accordingly, it scans an input rBPF program (*i.e.* a list of 64-bit bytecode instructions) and rejects it either cases violating *verifier_inv*:

- a source or destination register is greater than 10.
- the offset of a jump instruction is out of the instruction sequence bounds.
- the last instruction is not the *Exit* instruction (opcode 0x95).

Static verification of these simple properties allows for the interpreter to weave unnecessary dynamic checks as an asserted pre-condition. Moreover, our verifier adopts the same end-to-end verification method as the interpreter, Sec. 4.

The virtual machine state in CertrBPF-verifier is a strict subset of the interpreter's state: $st_v = \langle I, M \rangle$ consists of a sequence of instructions I and a memory M . The Coq theorem *verifier_well_formedness_and_safety* instantiates both Theorem 1 and Theorem 2 to state that the verifier 1) has no assumption (well-formedness); 2) never crashes (safety); 3) never modifies the VM state. In addition, the Coq theorem *verifier_imply_inv* states that *verifier* implies invariant *verifier_inv*.

```

Theorem verifier_well_formedness_and_safety :
  ∀ (st: verifier_state) (b: bool),
    verifier st = Some (b , st).

Theorem verifier_imply_inv :
  ∀ (st: verifier_state) (st': state)
    (Hinclude: st ⊂ st') (Hpre : verifier st = Some (true, st)),
    verifier_inv st'.

```

Considering that the verifier's proof and synthesis model are exactly the same, the simulation relation $R_v \subseteq st_v \times st_v$ required by Theorem 3 is *reflexivity*. Finally, CertrBPF-verifier reuses the *Clightlogic* to prove the simulation proof of its C implementation.

9 Evaluation: case study of RIOT’s femto-containers

In this section we evaluate CertrBPF in a concrete use case. We integrate CertrBPF as a drop-in replacement for the corresponding non-verified module optimized for size (vanilla-rBPF) in the IoT operating system RIOT to provide the expected femto-container functionalities [38].

Implementation Table 2 provides statistics on the complete specifications and proofs of CertrBPF. The proof model of the interpreter (Sec. 5) consists of 2.4k lines of Coq code and the corresponding isolation proof (Sec. 5.1) is more than 4.8k lines. The C-ready implementation (or synthesis model, Sec. 6) is approx. 3.2k lines long and the equivalence theorem is completed by 0.6k proof code. The final step (Sec. 7) includes 4.4k *Clightlogic* implementation and 10.8k translation validation proofs between the Gallina specification and the extracted Clight model. As for the CertrBPF verifier (Sec. 8), the proof and synthesis models sport 1.4k lines of Coq code. The corresponding property proofs spend more than 0.5k and the last simulation proof is about 8.3k long.

Table 2: Code statistics of CertrBPF

	Interpreter (locs)	Verifier (locs)
Proof Model	2445	1459
Properties	4885	547
Synthesis Model	3285	
Equivalence	635	
Clightlogic	4413	
Simulation	10820	8331
(Total)	26483	10337

Experimental evaluation setup. We carry out our measurements on a selected set of popular, commercial, off-the-shelf low-power IoT hardware, representative of modern 32-bit microcontroller architectures: Arm Cortex-M, Espressif ESP32, and RISC-V. Correspondingly, the boards we used are the Nordic nRF52840 development kit, the Espressif WROOM-32 board, and the Sipeed Longan Nano GD32VF103CBT6 development board. All code is compiled with GCC using size optimization enabled and the `-foptimize-sibling-calls` GCC option to reduce all tail-recursive calls and bound the stack size. This is critical to our isolation theorem as it relies on the implicit CompCert assumption that the stack cannot overflow.

Results. We first evaluate the memory footprint of the CertrBPF interpreter, compared to vanilla-rBPF in our use case. In terms of Flash, our measurements show that CertrBPF actually reduces the footprint by 37 % on RISC-V and by 25 % on ESP32, while incurring 6 % increase on Cortex-M. In terms of stack

requirements, CertrBPF reduces the footprint by 33 % on Cortex-M, by 22 % on RISC-V, and by 4 % on ESP32. The context memory, however, increases from 92 B to 144 B on all platforms.

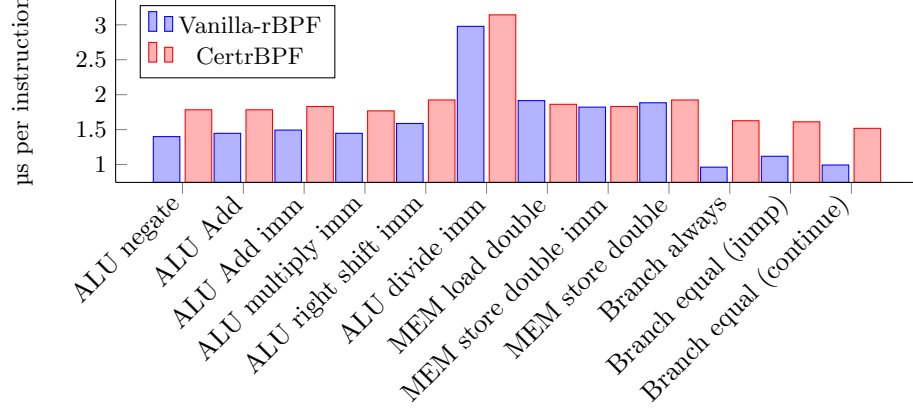


Fig. 5: Time per instructions on the Cortex-M4 platform

Next, we micro-benchmark the performance of fundamental operations: single instructions from the arithmetic logic unit (ALU), for memory access (MEM) and branch instructions, with a mix of register and immediate value for the operands, Figure 5. These results are averages over 1000 single identical instruction calls with a single return statement to make the application exit.

Finally, we benchmark the performance of actual IoT data processing, hosted in a femto-container with RIOT running on our selected hardware. In this use case, a sliding window average is performed within the femto-container, on available sensor data points. Figure 6 shows the performance we measured depending on the size of the window. We use this as blueprint for computation load scaling. To visualize the trade-off between runtime speed and flash size, we also ran sliding window benchmarks with forced function inlining on the certrBPF code. We observed that the resulting VM would then be 3208 B in size while averaging 1024 μs per execution versus 1341 μs, hinting at fine-tuning opportunities to increase speed by up to 30 %.




Fig. 6: Sliding window average on Cortex-M, ESP32, and RISC-V.

Key take-away. We observe that CertrBPF generally decreases the memory footprint. One reason is that calls to the RIOT API are currently not supported by CertrBPF. We observe, Figure 5, that the execution slow-down is most acute for Branch instructions, on Cortex-M. However, on all platforms (RISC-V, ESP32 and Cortex-M) our micro-benchmarks show that most instructions suffer minimal slow-down with CertrBPF compared to vanilla-rBPF. This behavior is also visible in our sensor data processing benchmark, Fig. 4, where the performances of CertrBPF and vanilla-rBPF are almost identical on ESP32 and RISC-V. All in all, the security gains brought by CertrBPF can be considered worth the performance loss we measured.

10 Related Works

Methodologies for systems and compilers verification The verification of compilers [16], static analyzers [14], and operating systems [15,10] have been the subjects of colossal development and verification due to the sheer code size of the artifacts at stake. These full-scale case studies gave rise to new strategies and methodologies to address the challenge of verifying large software. One key contribution in the development of the verified SeL4 micro-kernel, for instance, is the Cogent [32] approach: Cogent consists of a functional language with linear types to specify source programs and produces C code with Isabelle/HOL proof information. It provides a framework to prove that the extracted C code refines a high-level Isabelle/HOL functional correctness specification in the Isabelle/HOL proof assistant. Our method differs from co-specification in Cogent in that it is direct: it directly translates Coq specifications into C code and performs the end-to-end verification in Coq. CertiKOS [10] uses a multi-layered, refinement-based, and modular definition of a micro-kernel from its low-level memory model to its user-level interface and services. It is adopted in SeKVM [20], a layered Linux KVM hypervisor architecture for multiprocessor hardware. The CompCert project [16] adopted this “divide-and-conquer” strategy to decompose the verification of a full-scale ANSI C compiler into that of its successive transformations from source program to machine code, compositionally verifying each of the translation steps bisimilar. Its related static analyser, Verasco [14], employs static analysis of CompCert C code using a verified core abstract interpreter with composable abstract domains. Our problem statement is methodologically simpler: to build a safe and small VM that interprets rBPF virtual instructions on networked micro-controllers. We choose the radical approach of proof-oriented programming (à la Low*, Vale) to prove an rBPF interpreter embedded in Coq correct and to directly extract verified code from its definition.

Background on BPF and its verified implementations Mogul et al. [24] introduce a stack-based virtual machine to interpret packet filters into the BSD kernel that BPF extended to 32-bit instructions. BPF gained adoption in the Linux community and became eBPF (extended BPF), a virtual 64-bit RISC-like architectures. To our knowledge, verification of BPF runtime systems has mainly focused on

JIT translation for operation on micro-kernels. Myreen [26] verifies a JIT compiler from BPF to Intel architectures using the HOL4 proof assistant and an x86 compiler as TCB. Porncharoenwase et al. [31] use CompCert to extract an OCaml translator from BPF to assembly code, verified using the proof assistant Coq, using the OCaml runtime, an assembler, and a linker as TCB. Van Geffen et al. [36] present an optimized JIT compiler for Linux BPF with automated static analysis onboard, assuming offline verification using the Linux BPF verifier as TCB. For field deployment on networks of micro-controllers (IoT), all the above approaches would require a trusted, offline BPF verifier and, additionally, a secure upload protocol to sign verified scripts and perform authenticated uploads on target devices, which motivates our approach to use a fault-proof virtual machine instead.

Background on verified virtual machines Indeed, as Google’s Project Zero recently pointed out a series on JIT-compiler exploits, interest for isolating byte-code interpreters inside virtual machines has regained. Lochbihler [21] presents the verified implementation of a virtual machine modeling the semantics, memory model and byte-code semantics of Java, all by using the proof methodology of translation validation [30,16]. Desharnais and Brunthaler [6] propose the formal verification of an optimized and secure Javascript interpreter in Isabelle/HOL. Its proof methodology is based on concepts of bisimulation. The interpreter targets optimal security and run-time performance. To target MCU devices, our rBPF VM instead seeks optimal run-time memory footprint, to support the expected capability of dynamically running several isolated services on a small device with shared memory. Zhang et al. [39] presents a different and ambitious workflow using the deductive programming environment Why3 [8] to specify a virtual machine of Ethereum byte-code (EVM) and verify functional correctness of smart contracts against it. The EVM is extracted to OCaml binary code, yielding a TCB consisting of the OCaml runtime and the implementation of Ethereum’s protocols.

Background on converting Gallina programs into executables Just as the proof-oriented approach advocated by dependently-typed functional languages like F* mentioned in Sec. 2, there are various alternatives to ∂x for extracting executables from Gallina programs. To begin with, Coq comes with a builtin extraction mechanism [19] that generates OCaml, Haskell or Scheme. This path has a rather large TCB (Coq extraction and a compiler). CertiCoq [1] is an ongoing project aiming at generating CompCert C code from Gallina using a specific IR and several passes. Once this effort is completed, it will allow one to rely on a small TCB. Euf [25] is another tool to compile Gallina to C. It considers a carefully chosen subset of Gallina to tackle the tricky issue of verifying the reflection of Gallina into an AST. Both CertiCoq and Euf, however, require a garbage collector and define how Coq inductives are represented at runtime. Codegen [34] converts Gallina to C with the goal of maximizing performance by, *e.g.* allowing the user to control how Coq values are represented at runtime. Rupicola [29] considers an original and promising approach which regards a compiler as a

partial decision procedure: it consists of a proof search procedure, which may fail, or else exhibit a target program in bedrock2 (a C-like low-level language AST of Coq) with a proof of equivalence. It has, at present, only been tested for small algorithms. We chose to use ∂x for its simplicity and because it doesn't increase our TCB. It shares with codegen the capability to configure the representation of values. Unlike codegen, it produces C code that is structurally identical to source code. This direct and traceable translation simplifies the verification of generated code w.r.t. source programs, and facilitates source program optimisations.

11 Conclusion and Future Works

This paper proposes a refinement methodology to directly derive a verified C implementation of rBPF, the implementation of BPF hosted by the RIOT operating system, from a Gallina specification in Coq. All the refinement steps are mechanically verified using the Coq proof assistant to minimize the TCB. We prove our rBPF virtual machine to isolate software faults and not to produce runtime errors. Performances are at par with the vanilla rBPF implementation in RIOT which, yet, uses GCC's `label-as-values` extensions. As these extensions cannot be expressed by the CompCert semantics, the optimisations they provide are out of reach to CertrBPF.

Our future works aim at instantiating our proof method to Wasm, or to that of a (fault-isolating) JIT compiler, one challenge being that Linux's approach of using a verifier will not be feasible on resource-constrained devices, and another being that certain operations might only be expressible in assembly code. This calls for further studies on ways to substantially improve the efficiency of our VM.

References

1. Anand, A., Appel, A.W., Morrisett, G., Paraskevopoulou, Z., Pollack, R., Bélanger, O.S., Sozeau, M., Weaver, M.Z.: Certicoq : A verified compiler for coq. In: CoqPL workshop. Paris, France (2017)
2. Appel, A.W., Dockins, R., Hobor, A., Beringer, L., Dodds, J., Stewart, G., Blazy, S., Leroy, X.: Program logics for certified compilers. Cambridge University Press (2014)
3. Baccelli, E., Gündoğan, C., Hahm, O., Kietzmann, P., Lenders, M.S., Petersen, H., Schleiser, K., Schmidt, T.C., Wählich, M.: RIOT: An open source operating system for low-end embedded devices in the IoT. *IEEE Internet of Things Journal* **5**(6), 4428–4440 (2018)
4. Bertot, Y., Castéran, P.: Interactive theorem proving and program development: Coq’Art: the calculus of inductive constructions. Springer Science & Business Media, Berlin, Heidelberg (2013)
5. Costan, V., Lebedev, I., Devadas, S.: Sanctum: Minimal hardware extensions for strong software isolation. In: 25th USENIX Security Symposium (USENIX Security 16). pp. 857–874. USENIX Association, Austin, TX (Aug 2016)
6. Desharnais, M., Brunthaler, S.: Towards efficient and verified virtual machines for dynamic languages. In: Proceedings of the 10th ACM SIGPLAN International Conference on Certified Programs and Proofs. p. 61–75. CPP 2021, ACM, New York, NY, USA (2021)
7. Dunchev, C., Guidi, F., Coen, C.S., Tassi, E.: ELPI: Fast, embeddable, λ Prolog interpreter. In: Davis, M., Fehnker, A., McIver, A., Voronkov, A. (eds.) *Logic for Programming, Artificial Intelligence, and Reasoning - 20th International Conference, LPAR-20 2015, Suva, Fiji, November 24–28, 2015, Proceedings. Lecture Notes in Computer Science*, vol. 9450, pp. 460–468. Springer, Suva, Fiji (2015). https://doi.org/10.1007/978-3-662-48899-7_32, https://doi.org/10.1007/978-3-662-48899-7_32
8. Filliâtre, J.C., Paskevich, A.: Why3—where programs meet provers. In: European symposium on programming, pp. 125–128. Springer Berlin Heidelberg, Rome, Italy (2013)
9. Fleming, M.: A Thorough Introduction to eBPF. *Linux Weekly News* (2017)
10. Gu, L., Vaynberg, A., Ford, B., Shao, Z., Costanzo, D.: CertiKOS: a certified kernel for secure cloud computing. In: Proceedings of the Second Asia-Pacific Workshop on Systems - APSys ’11. p. 1. ACM Press, Shanghai, China (2011). <https://doi.org/10.1145/2103799.2103803>, <http://dl.acm.org/citation.cfm?doid=2103799.2103803>
11. Haas, A., Rossberg, A., Schuff, D.L., Titzer, B.L., Holman, M., Gohman, D., Wagner, L., Zakai, A., Bastien, J.: Bringing the web up to speed with webassembly. In: Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation. p. 185–200. PLDI 2017, Association for Computing Machinery, New York, NY, USA (2017). <https://doi.org/10.1145/3062341.3062363>, <https://doi.org/10.1145/3062341.3062363>
12. Hahm, O., Baccelli, E., Petersen, H., Tsiftes, N.: Operating Systems for Low-End Devices in the Internet of Things: a Survey. *IEEE internet of things journal* **3**(5), 720–734 (Oct 2016)
13. Jomaa, N., Torrini, P., Nowak, D., Grimaud, G., Hym, S.: Proof-oriented design of a separation kernel with minimal trusted computing base. In: 18th International Workshop on Automated Verification of Critical Systems (AVOCS 2018). vol. 76.

- Electronic Communications of the EASST Open Access Journal, Oxford, United Kingdom (Jul 2018)
14. Jourdan, J.H., Laporte, V., Blazy, S., Leroy, X., Pichardie, D.: A formally-verified C static analyzer. In: Proceedings of the 42nd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages. pp. 247–259. ACM, Mumbai India (Jan 2015). <https://doi.org/10.1145/2676726.2676966>, <https://dl.acm.org/doi/10.1145/2676726.2676966>
 15. Klein, G., Norrish, M., Sewell, T., Tuch, H., Winwood, S., Elphinstone, K., Heiser, G., Andronick, J., Cock, D., Derrin, P., Elkaduwe, D., Engelhardt, K., Kolanski, R.: seL4: formal verification of an OS kernel. In: Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles - SOSP '09. p. 207. ACM Press, Big Sky, Montana, USA (2009). <https://doi.org/10.1145/1629575.1629596>, <http://portal.acm.org/citation.cfm?doid=1629575.1629596>
 16. Leroy, X.: Formal verification of a realistic compiler. Communications of the ACM **52**(7), 107–115 (Jul 2009). <https://doi.org/10.1145/1538788.1538814>, <https://dl.acm.org/doi/10.1145/1538788.1538814>
 17. Leroy, X., Appel, A.W., Blazy, S., Stewart, G.: The CompCert Memory Model, Version 2. Research Report RR-7987, INRIA (Jun 2012)
 18. Leroy, X., Blazy, S.: Formal verification of a C-like memory model and its uses for verifying program transformations. Journal of Automated Reasoning **41**(1), 1–31 (2008)
 19. Letouzey, P.: A new extraction for coq. In: International Workshop on Types for Proofs and Programs. pp. 200–219. Springer, Berg en Dal, The Netherlands (2002)
 20. Li, S.W., Li, X., Gu, R., Nieh, J., Hui, J.Z.: Formally verified memory protection for a commodity multiprocessor hypervisor. In: 30th USENIX Security Symposium (USENIX Security 21). pp. 3953–3970. USENIX Association, Virtual Events (Aug 2021)
 21. Lochbihler, A.: A Machine-Checked, Type-Safe Model of Java Concurrency : Language, Virtual Machine, Memory Model, and Verified Compiler. Ph.D. thesis, Karlsruhe Institute of Technology (2012). <https://doi.org/10.5445/KSP/1000028867>
 22. Madhavapeddy, A., Mortier, R., Rotsos, C., Scott, D., Singh, B., Gazagnaire, T., Smith, S., Hand, S., Crowcroft, J.: Unikernels: Library operating systems for the cloud. In: Proceedings of the Eighteenth International Conference on Architectural Support for Programming Languages and Operating Systems. p. 461–472. ASPLOS '13, Association for Computing Machinery, New York, NY, USA (2013). <https://doi.org/10.1145/2451116.2451167>, <https://doi.org/10.1145/2451116.2451167>
 23. McCanne, S., Jacobson, V.: The bsd packet filter: A new architecture for user-level packet capture. In: Proceedings of the Usenix Winter 1993 Technical Conference, San Diego, California, USA, January 1993. vol. 46, pp. 259–270. USENIX Association, San Diego, California, USA (1993)
 24. Mogul, J., Rashid, R., Accetta, M.: The packer filter: an efficient mechanism for user-level network code. ACM SIGOPS Operating Systems Review **21**(5), 39–51 (1987)
 25. Mullen, E., Pernsteiner, S., Wilcox, J.R., Tatlock, Z., Grossman, D.: Œuf: minimizing the coq extraction TCB. In: Andronick, J., Felty, A.P. (eds.) Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8-9, 2018. pp. 172–185. ACM, Los Angeles, CA, USA (2018). <https://doi.org/10.1145/3167089>

26. Myreen, M.O.: Verified just-in-time compiler on x86. In: Proceedings of the 37th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages. pp. 107–118. ACM New York, NY, USA, Madrid, Spain (2010)
27. Nelson, L., Geffen, J.V., Torlak, E., Wang, X.: Specification and verification in the field: Applying formal methods to BPF just-in-time compilers in the linux kernel. In: 14th USENIX Symposium on Operating Systems Design and Implementation (OSDI 20). pp. 41–61. USENIX Association, Virtual Events (Nov 2020), <https://www.usenix.org/conference/osdi20/presentation/nelson>
28. Noorman, J., Agten, P., Daniels, W., Strackx, R., Herrewewege, A.V., Huygens, C., Preneel, B., Verbauwhede, I., Piessens, F.: Sancus: Low-cost trustworthy extensible networked devices with a zero-software trusted computing base. In: 22nd USENIX Security Symposium (USENIX Security 13). pp. 479–498. USENIX Association, Washington, D.C. (Aug 2013), <https://www.usenix.org/conference/usenixsecurity13/technical-sessions/presentation/noorman>
29. Pit-Claudel, C., Philipoom, J., Jamner, D., Erbsen, A., Chlipala, A.: Relational compilation for performance-critical applications. In: Proceedings of the 43th ACM SIGPLAN Conference on Programming Language Design and Implementation. PLDI 2022, Association for Computing Machinery, New York, NY, USA (2022)
30. Pnueli, A., Siegel, M., Singerman, E.: Translation validation. In: Goos, G., Hartmanis, J., van Leeuwen, J., Steffen, B. (eds.) Tools and Algorithms for the Construction and Analysis of Systems, vol. 1384, pp. 151–166. Springer Berlin Heidelberg, Berlin, Heidelberg (1998). <https://doi.org/10.1007/BFb0054170>, <http://link.springer.com/10.1007/BFb0054170>, series Title: Lecture Notes in Computer Science
31. Porncharoenwase, S., Bornholt, J., Torlak, E.: Fixing code that explodes under symbolic evaluation. In: International Conference on Verification, Model Checking, and Abstract Interpretation. pp. 44–67. Springer, New Orleans, LA, USA (2020)
32. Rizkallah, C., Lim, J., Nagashima, Y., Sewell, T., Chen, Z., O’Connor, L., Murray, T., Keller, G., Klein, G.: A framework for the automatic formal verification of refinement from cogent to c. In: International Conference on Interactive Theorem Proving. pp. 323–340. Springer, Nancy, France (2016)
33. Ševčík, J., Vafeiadis, V., Zappa Nardelli, F., Jagannathan, S., Sewell, P.: Comperttso: A verified compiler for relaxed-memory concurrency. *Journal of the ACM (JACM)* **60**(3), 1–50 (2013)
34. Tanaka, A.: Coq to C translation with partial evaluation. In: Lindley, S., Mogensen, T.Æ. (eds.) Proceedings of the 2021 ACM SIGPLAN Workshop on Partial Evaluation and Program Manipulation, PEPM@POPL 2021, Virtual Event, Denmark, January 18–19, 2021. pp. 14–31. ACM, New York, NY, United States (2021). <https://doi.org/10.1145/3441296.3441394>
35. Tassi, E.: Coq-Elpi, Coq plugin embedding Elpi (2021), <https://github.com/LPCIC/coq-elpi>
36. Van Geffen, J., Nelson, L., Dillig, I., Wang, X., Torlak, E.: Synthesizing jit compilers for in-kernel dsls. In: Lahiri, S.K., Wang, C. (eds.) Computer Aided Verification. pp. 564–586. Springer International Publishing, Cham (2020)
37. Zandberg, K., Baccelli, E.: Minimal virtual machines on IoT microcontrollers: The case of Berkeley Packet Filters with rBPF. In: 2020 9th IFIP International Conference on Performance Evaluation and Modeling in Wireless Networks (PEMWN). pp. 1–6. IEEE, Berlin / Virtual, Germany (2020)
38. Zandberg, K., Baccelli, E.: Femto-Containers: DevOps on Microcontrollers with Lightweight Virtualization & Isolation for IoT Software Modules (Nov 2021), <https://hal.inria.fr/hal-03263164>, preprint

39. Zhang, X., Li, Y., Sun, M.: Towards a formally verified EVM in production environment. In: Bliudze, S., Bocchi, L. (eds.) *Coordination Models and Languages*. pp. 341–349. Springer International Publishing, Cham (2020)

Appendix: source code, proofs and benchmarks

The source code and proofs of our virtual machine, its generated code and benchmark data are available to the CAV'22 evaluation committee on the anonymized repository: <https://anonymous.4open.science/r/AnonymousCAV22-59DC>.